

1. What problem does an output parser solve when working with LLMs? Explain why unstructured text can be difficult for software applications to process.
2. Explain the complete flow of a StructuredOutputParser, starting from defining a ResponseSchema and ending with the final structured Python output.
3. What is the difference between JsonOutputParser and PydanticOutputParser? Explain how their final outputs differ.
4. What is the difference between PydanticOutputParser and with_structured_output()? When would you choose one over the other?
5. Why might an output parser fail even after format instructions have been provided to the LLM? Discuss at least two possible reasons.